ADVANCED COMMUNITY DATA ANALYSIS USING VEGAN

Gavin L. Simpson ESA 2015 · Aug 9th 2015



CANONICAL CORRESPONDENCE ANALYSIS

CCA is the constrained form of CA; fitted using cca().

Two interfaces for specifying models

- basic; cca1 <- cca(X = varespec, Y = varechem)</pre>
- formula; cca1 <- cca(varespec ~ ., data = varechem)</pre>

Formula interface is the more powerful — recommended

CANONICAL CORRESPONDENCE ANALYSIS

```
cca1 <- cca(varespec ~ ., data = varechem)</pre>
cca1
Call: cca(formula = varespec ~ N + P + K + Ca + Mg + S + Al + Fe +
Mn + Zn + Mo + Baresoil + Humdepth + pH, data = varechem)
             Inertia Proportion Rank
Total
            2.0832 1.0000
Constrained 1.4415 0.6920 14
Unconstrained 0.6417 0.3080
Inertia is mean squared contingency coefficient
Eigenvalues for constrained axes:
 CCA1 CCA2 CCA3 CCA4 CCA5 CCA6 CCA7 CCA8 CCA9 CCA10
0 4389 0 2918 0 1628 0 1421 0 1180 0 0890 0 0703 0 0584 0 0311 0 0133
CCA11 CCA12 CCA13 CCA14
0.0084 0.0065 0.0062 0.0047
Eigenvalues for unconstrained axes:
   CA1
           CA2
                   CA3
                          CA4
                                 CA5
                                         CA6
0.19776 0.14193 0.10117 0.07079 0.05330 0.03330 0.01887 0.01510 0.00949
```

REDUNDANCY ANALYSIS

RDA is the constrained form of PCA; fitted using rda().

```
rda1 <- rda(varespec ~ .. data = varechem)
rda1
Call: rda(formula = varespec ~ N + P + K + Ca + Mg + S + Al + Fe +
Mn + Zn + Mo + Baresoil + Humdepth + pH. data = varechem)
               Inertia Proportion Rank
Total
            1825.7000
                         1.0000
Constrained 1459,9000
                       0.7997 14
Unconstrained 365,8000
                          0.2003 9
Inertia is variance
Eigenvalues for constrained axes:
RDA1 RDA2 RDA3 RDA4 RDA5 RDA6
                                  RDA7 RDA8 RDA9 RDA10 RDA11 RDA12
820.1 399.3 102.6 47.6 26.8 24.0 19.1 10.2 4.4 2.3 1.5 0.9
RDA13 RDA14
 0.7 0.3
Eigenvalues for unconstrained axes:
  PC1
         PC2
               PC3
                      PC4
                             PC5
                                   PC6
                                                       PC9
186.19 88.46 38.19 18.40 12.84 10.55
                                        5.52 4.52 1.09
```

THE cca.object

- · Objects of class "cca" are complex with many components
- Entire class described in ?cca.object
- Depending on what analysis performed some components may be **NULL**
- Used for (C)CA, PCA, RDA, and CAP (capscale())

THE cca.object

cca1 has a large number of components

- **\$call** how the function was called
- \$grand.total in (C)CA sum of 'rowsum}
- **\$rowsum** the row sums
- **\$colsum** the column sums
- **\$tot.chi** total inertia, sum of Eigenvalues
- *pCCA Conditioned (partialled out) components
- **\$CCA** Constrained components
- \$CA Unconstrained components
- **\$method** Ordination method used
- **\$inertia** Description of what inertia is

THE cca.object

Depending on how one called cca() etc some of these components will be NULL

\$pCCA is only filled in if a *partial* constrained ordination fitted

rda() returns objects with classes "rda" and "cca", but in most cases those objects
work like those of class "cca"

The Eigenvalues and axis scores are now spread about the **\$CA** and **\$CCA** components (also **\$pCCA** if a *partial* CCA)

Thankfully we can use extractor functions to get at such things

EIGENVALUES

Use eigenvals() to extract Eigenvalues from a fitted ordination object

```
eigenvals(cca1)
```

```
CCA1
               CCA2
                         CCA3
                                   CCA4
                                              CCA5
                                                        CCA6
                                                                  CCA7
0.4388704 0.2917753 0.1628465 0.1421302 0.1179519 0.0890291 0.0702945
     CCA8
               CCA9
                        CCA10
                                  CCA11
                                             CCA12
                                                       CCA13
                                                                 CCA14
0.0583592 0.0311408 0.0132944 0.0083644 0.0065385 0.0061563 0.0047332
      CA1
                CA2
                          CA3
                                    CA4
                                               CA5
                                                         CA6
                                                                   CA7
0.1977645 0.1419256 0.1011741 0.0707868 0.0533034 0.0332994 0.0188676
      CA8
                CA9
0.0151044 0.0094876
```

Ç

EXTRACTING AXIS SCORES

To extract a range of scores from a fitted ordination use scores()

- takes an ordination object as the first argument
- choices which axes? Defaults to c(1,2)
- display which type(s) of scores to return
 - "sites" or "wa": scores for samples in response matrix
 - "species": scores for variables/columns in response
 - "lc": linear combination site scores
 - "bp": biplot scores (coords of arrow tip)
 - "cn": centroid scores (coords of factor centroids)

EXTRACTING AXIS SCORES

```
str(scores(cca1, choices = 1:4, display = c("species", "sites")), max = 1)
list of 2
 $ species: num [1:44, 1:4] 0.0753 -0.1813 -1.0535 -1.2774 -0.1526 ...
  ..- attr(*. "dimnames")=List of 2
 $ sites : num [1:24, 1:4] 0.178 -0.97 -1.28 -1.501 -0.598 ...
  ..- attr(*. "dimnames")=List of 2
head(scores(cca1, choices = 1:2, display = "sites"))
        CCA1
               CCA2
18 0.1784733 -1.0598842
15 -0.9702382 -0.1971387
24 -1.2798478 0.4764498
27 -1.5009195 0.6521559
23 -0.5980933 -0.1840362
19 -0.1102881 0.7143142
```

SCALINGS...

When we draw the results of many ordinations we display 2 or more sets of data

Can't display all of these and maintain relationships between the scores

Solution scale one set of scores relative to the other via the scaling argument

- · scaling = 1 Focus on sites, scale site scores by λ_i
- \cdot scaling = 2 Focus on species, scale species scores by λ_i
- \cdot scaling = 3 Symmetric scaling, scale both scores by $\sqrt{\lambda_i}$
- \cdot scaling = -1 As above, but
- · scaling = -2 For cca() multiply results by $\sqrt{(1/(1-\lambda_i))}$
- scaling = -3 this is Hill's scaling
- scaling < 0 For rda() divide species scores by species' σ
- scaling = 0 raw scores

```
scores(cca1, choices = 1:2, display = "species", scaling = 3)
```

PARTIAL CONSTRAINED ORDINATIONS

Partial constrained ordinations remove the effect of one or more variables then fit model of interest

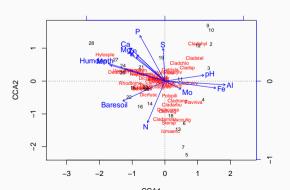
Argument Z is used for a data frame of variables to partial out

Or with the formula interface use the Condition() function

TRIPLOTS

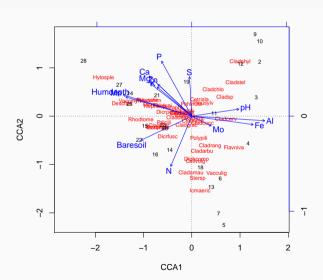
Triplots will generally produce a mess; we can really only display a couple of bits approximately anyway Trying to cram three things in is a recipe for a mess... but we can do it

plot(cca1)



TRIPLOTS

plot(cca1)



BUILDING CONSTRAINED ORDINATION MODELS

If we don't want to think it's easy to fit a poor model with many constraints

That's what we just did with cca1 and rda1

Remember, CCA and RDA are *just regression methods* — everything you know about regression applies here

A better approach is to *think* about the important variables and include only those

The formula interface allows you to create interaction or quadratic terms easily (though be careful with latter)

It also handles factor or class constraints automatically unlike the basic interface

BUILDING CONSTRAINED ORDINATION MODELS

```
vare.cca <- cca(varespec ~ Al + P*(K + Baresoil), data = varechem)
vare.cca
Call: cca(formula = varespec \sim Al + P * (K + Baresoil). data =
varechem)
            Inertia Proportion Rank
Total
         2.083 1.000
Constrained 1.046 0.502 6
Unconstrained 1.038 0.498 17
Inertia is mean squared contingency coefficient
Eigenvalues for constrained axes:
 CCA1 CCA2 CCA3 CCA4 CCA5 CCA6
0.3756 0.2342 0.1407 0.1323 0.1068 0.0561
Eigenvalues for unconstrained axes:
   CA1
           CA2
                  CA3 CA4 CA5 CA6 CA7
                                                       CA8
0.27577 0.15411 0.13536 0.11803 0.08887 0.05511 0.04919 0.03781
(Showed only 8 of all 17 unconstrained eigenvalues)
```

BUILDING CONSTRAINED ORDINATION MODELS

For CCA we have little choice but to do

- 1. Fit well-chosen set of candidate models & compare, or
- 2. Fit a full model of well-chosen variables & then do stepwise selection

But automatic approaches to model building should be used cautiously!

The standard step() function can be used as vegan provides two helper methods, deviance() and extractAIC(), used by step()

Vegan also provides methods for class "cca" for add1() and drop1()

VARIANCE INFLATION FACTORS

Linear dependencies between constraints can be investigated via the variance inflation factor or VIF

VIF is a measure of how much the variance of $\hat{\beta}_j$ is inflated by presence of other covariates

Lots of rules of thumb

- VIF >= 20 indicates strong collinearity in constraints
- VIF >= 10 potnetially of concern & should be looked at

Computed via vif.cca()

STEPWISE SELECTION IN CCA

step() uses AIC which is a fudge for RDA/CCA. Alternatively use function ordistep()

- 1. Define an upper and lower model scope, say the full model and the null model
- 2. To step from the lower scope or null model we use

```
upr <- cca(varespec ~ ., data = varechem)
lwr <- cca(varespec ~ 1, data = varechem)
set.seed(1)
mods <- ordistep(lwr, scope = formula(upr), trace = 0)</pre>
```

trace = 0 is used her to turn off printing of progress

Permutation tests are used (more on these later); the theory for an AIC for ordination is somewhat loose

STEPWISE SELECTION IN CCA

The object returned by **step()** is a standard **"cca"** object with an extra component **\$anova**

The **\$anova** component contains a summary of the steps involved in automatic model building

```
mods
Call: cca(formula = varespec ~ Al + P + K, data = varechem)
             Inertia Proportion Rank
Total
              2.0832
                         1 0000
Constrained
              0.6441
                         0.3092
Unconstrained 1,4391
                      0.6908
Inertia is mean squared contingency coefficient
Eigenvalues for constrained axes:
 CCA1 CCA2 CCA3
0 3616 0 1700 0 1126
Eigenvalues for unconstrained axes:
          CA2
                CA3
                        CA4
                              CA5
                                     CA6
A 3500 A 2201 A 1851 A 1551 A 1351 A 1003 A 0773 A 0537
(Showed only 8 of all 20 unconstrained eigenvalues)
```

STEPWISE SELECTION IN CCA

The **\$anova** component contains a summary of the steps involved in automatic model building

mods\$anova

```
Df AIC F Pr(>F)
+ Al 1 128.61 3.6749 0.005 **
+ P 1 127.91 2.5001 0.005 **
+ K 1 127.44 2.1688 0.035 *
---
Signif. codes: 0 '***' 0.001 '**' 0.05 '.' 0.1 ' ' 1
```

Step-wise model selection is fairly fragile; if we start from the full model we won't end up with the same final model

```
mods2 <- step(upr, scope = list(lower = formula(lwr), upper = formula(upr)), trace = 0,</pre>
             test = "perm")
mods2
Call: cca(formula = varespec ~ P + K + Mg + S + Mn + Mo + Baresoil
+ Humdepth, data = varechem)
             Inertia Proportion Rank
              2.0832
                         1.0000
Total
Constrained
              1.1165
                      0.5360
Unconstrained A.9667
                      0.4640
                                  15
Inertia is mean squared contingency coefficient
Eigenvalues for constrained axes:
  CCA1 CCA2
               CCA3 CCA4 CCA5 CCA6 CCA7 CCA8
0.4007 0.2488 0.1488 0.1266 0.0875 0.0661 0.0250 0.0130
Eigenvalues for unconstrained axes:
                           CA4
                                   CA5
                                                   CA7
    CA1
            CA2
                   CA3
                                           CAG
                                                           CAR
                                                                   CAG
0.25821 0.18813 0.11927 0.10204 0.08791 0.06085 0.04461 0.02782 0.02691
   CA10
          CA11
                  CA12
                          CA13
                                  CA14
                                          CA15
0.01646 0.01364 0.00823 0.00655 0.00365 0.00238
```

ADJUSTED $\it r^2$ FOR $\it LINEAR$ MODELS

As with ordinary \mathbb{R}^2 , that of an RDA is biased for the same reasons as for a linear regression

- \cdot adding a variable to constraints will increase R^2
- \cdot the larger the number of constraints in the model the larger \mathbb{R}^2 is due to random correlations

Can attempt to account for this bias via an adjusted \mathbb{R}^2 measure

$$R_{adj}^2 = 1 - \frac{n-1}{n-m-1}(1-R^2)$$

- \cdot n is number of samples m is number of constraints (model degrees of freedom)
- Can be used up to $\sim M > n/2$ before becomes too conservative
- Can be negative
- Compute using RsquareAdj()

STEPWISE SELECTION VIA ADJUSTED $\it r^2$

The problems with stepwise selection in regression models are myriad. Affects RDA, CCA, etc as well

Blanchet, Legendre, and Borcard (2008) proposed a two-step solution for models where R^2_{adj} makes sense

- · Global test of all constraints
 - · Proceed only if this test is significant
 - Helps prevent inflation of overall type I error
- · Proceed with forward selection, but with two stopping rules
 - \cdot Usual significance threshold lpha
 - The global R^2_{adj}
 - · Stop if next candidate model is non-significant or if R^2_{adj} exceeds the global R^2_{adj}

Available in ordiR2step()



PERMUTATION TESTS IN VEGAN

RDA has lots of theory behind it, CCA not as much. However, ecological/environmental data invariably violate what little theory we have

Instead we use permutation tests to assess the importance of fitted models — the data are shuffled in some way and the model refitted to derive a Null distribution under some hypothesis of $no\ effect$

PERMUTATION TESTS IN VEGAN

What is shuffled and how is of paramount importance for the test to be valid

- · No conditioning (partial) variables then rows of the species data are permuted
- With conditioning variables, two options are available, both of which permute residuals from model fits
 - The *full model* uses residuals from model $Y = X + Z + \varepsilon$
 - The reduced model uses residuals from model $Y = X + Z + \varepsilon$
- In vegan which is used can be set via argument model with "direct", "full", and "reduced" respectively
- In current vegan option method = "full" is disabled

PERMUTATION TESTS IN VEGAN

A test statistic is required, computed for observed model & each permuted model vegan uses a pseudo-F statistic

$$F = \frac{\chi^2_{model}/df_{model}}{\chi^2_{resid}/df_{resid}}$$

Evaluate whether \it{F} is unusually large relative to the null (permutation) distribution of \it{F}

PERMUTATION TESTS IN VEGAN: anova()

set.seed(42)

- The main user function is the anova() method
- It is an interface to the lower-level function permutest.cca()
- At its most simplest, the anova() method tests whether the "model" as a whole is significant

$$F = \frac{1.4415/14}{0.6417/9} = 1.4441$$

PERMUTATION TESTS IN VEGAN: anova()

perm

- anova() will continue permuting only as long as it is uncertain whether the the p-value is above or below the chosen threshold (say p=0.05)
- If the function is sure the permuted p is above the threshold ${\tt anova()}$ may return after only a few hundred permutations
- In other cases many hundreds or thousands of permutations may be required to say whether the model is above or below the threshold
- · In the example, 1599 permutations were required

PERMUTATION TESTS IN VEGAN: anova()

· anova.cca() has a number of arguments

```
args(anova.cca)

function (object, ..., permutations = how(nperm = 999), by = NULL,
   model = c("reduced", "direct", "full"), parallel = getOption("mc.cores"),
   strata = NULL, cutoff = 1, scope = NULL)
NULL
```

- alpha is the desired p value threshold (Type I error rate)
- beta is the Type II error rate
- Permuting stops if the result is different from alpha for the given beta
- This is evaluated every step permutations
- perm.max sets a limit on the number of permutations
- by determines what is tested; the default is to test the model
- More direct control can be achieved via permutest.cca()

TYPES OF PERMUTATION TEST IN VEGAN

A number of types of test can be envisaged

- Testing the overall significance of the model
- Testing constrained (canonical) axes
- Testing individual model terms sequentially
- The marginal effect of a single variable

The first is the default in anova()

The other three can be selected via the argument method

PERMUTATION TESTS | TESTING CANONICAL AXES

- The constrained (canonical) axes can be individually tests by specifying by = "axis"
- The first axis is tested in terms of variance explained compared to residual variance
- The second axis is tested after partialling out the first axis... and so on

```
set.seed(1)
anova(mods. by = "axis")
Permutation test for cca under reduced model
Marginal tests for axes
Permutation: free
Number of permutations: 999
Model: cca(formula = varespec ~ Al + P + K, data = varechem)
        Df ChiSquare
                          F Pr(>F)
CCA1
             0.36156 5.0249 0.001 ***
     1 0.16996 2.3621 0.011 *
CCA2
      1 0.11262 1.5651 0.124
CCA3
Residual 20 1.43906
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

PERMUTATION TESTS | TESTING TERMS SEQUENTIALLY

- The individual terms in the model can be tested using by = "terms"
- The terms are assessed in the order they were specified in the model, sequentially from first to last
- Test is of the additional variance explained by adding the kth variable to the model
- Ordering of the terms will affect the results

PERMUTATION TESTS | TESTING TERMS MARGINAL EFFECTS

- The marginal effect of a model term can be assessed using by = "margin"
- The marginal *effect* is the effect of a particular term when all other model terms are included in the model

```
set.seed(10)
anova(mods, by = "margin")
Permutation test for cca under reduced model
Marginal effects of terms
Permutation: free
Number of permutations: 999
Model: cca(formula = varespec ~ Al + P + K. data = varechem)
        Df ChiSquare
                          F Pr(>F)
Δ1
         1 0.31184 4.3340 0.001 ***
        1 0.16810 2.3362 0.012 *
        1 0.15605 2.1688 0.025 *
Residual 20 1.43906
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Example & data taken from Leps & Smilauer, Case Study 2

Spring fen meadow vegetation in westernmost Carpathian mountains

```
## load vegan
library("vegan")

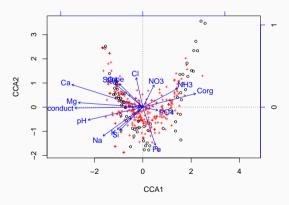
## load the data
spp <- read.csv("data/meadow-spp.csv", header = TRUE, row.names = 1)
env <- read.csv("data/meadow-env.csv", header = TRUE, row.names = 1)</pre>
```

 $m1 < -cca(spp \sim ... data = env)$

CCA a reasonable starting point as the gradient is long here (check with **decorana()** if you want)

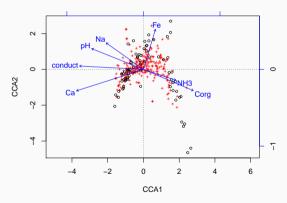
```
set.seed(32)
anova(m1)
Permutation test for cca under reduced model
Permutation: free
Number of permutations: 999
Model: cca(formula = spp ~ Ca + Mg + Fe + K + Na + Si + SO4 + PO4 + NO3 + NH3 + Cl + Corg + pH + conduct + slope.
        Df ChiSquare F Pr(>F)
Model 15 1.5597 1.497 0.001 ***
Residual 54 3.7509
Signif. codes: 0 '***' 0.001 '**' 0.05 '.' 0.1 ' ' 1
```

plot(m1)



```
set.seed(67)
lwr < -cca(spp ~ 1. data = env)
m2 <- ordistep(lwr. scope = formula(m1), trace = FALSE)
m2
Call: cca(formula = spp ~ Ca + conduct + Corg + Na + NH3 + Fe +
pH, data = env)
             Inertia Proportion Rank
Total
              5.3110
                     1 0000
Constrained
              0.9900
                     0.1864
Unconstrained 4.3210 0.8136 62
Inertia is mean squared contingency coefficient
Eigenvalues for constrained axes:
 CCA1 CCA2 CCA3 CCA4 CCA5 CCA6 CCA7
0.4268 0.1447 0.1116 0.0936 0.0760 0.0719 0.0652
Eigenvalues for unconstrained axes:
    CA1
           CA2
                   CA3
                          CA4
                                  CA5
                                          CA6
                                                 CA7
0.27251 0.19518 0.16703 0.14993 0.14606 0.14168 0.13292 0.12154
(Showed only 8 of all 62 unconstrained eigenvalues)
```

plot(m2)



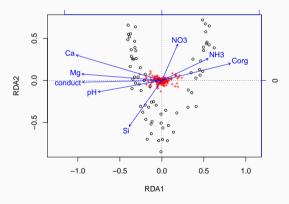
m2\$anova

```
Df AIC F Pr(>F)
+ Ca 1 453.14 4.7893 0.005 **
+ conduct 1 453.29 1.7915 0.005 **
+ Corg 1 453.61 1.6011 0.005 **
+ Na 1 453.93 1.5827 0.005 **
+ NH3 1 454.36 1.4507 0.020 *
+ Fe 1 454.89 1.3386 0.015 *
+ pH 1 455.46 1.2756 0.015 *
---
Signif. codes: 0 '***' 0.001 '**' 0.05 '.' 0.1 ' ' 1
```

Alternative is RDA with a transformation

```
spph <- decostand(spp, method = "hellinger")
m3 <- rda(spph ~ ., data = env)
lwr <- rda(spph ~ 1, data = env)
m4 <- ordistep(lwr, scope = formula(m3), trace = FALSE)</pre>
```

plot(m4)



Stepwise using ${\cal R}^2_{adj}$

```
m5 <- ordiR2step(lwr, scope = formula(m3), trace = FALSE)
m5$anova
```

DIAGNOSTICS FOR CONSTRAINED ORDINATIONS

vegan provides a series of diagnostics to help assess the model fit

- goodness()
- · inertcomp()
- · spenvcor()
- · intersetcor()
- · vif.caa()

DIAGNOSTICS FOR CONSTRAINED ORDINATIONS | GOODNESS OF FIT

goodness() computes two goodness of fit statistics for species or sites
Which goodness of fit measure is returned is controlled by argument statistic

- statistic = "explained (default) gives the cumulative proportion of variance explained by each axis
- statistic = "distance" gives the residual distance between the "fitted" location in constrained ordination space and the location in the full dimensional space

DIAGNOSTICS FOR CONSTRAINED ORDINATIONS | GOODNESS OF FIT

head(goodness(mods))

```
    CCA1
    CCA2
    CCA3

    Callvulg
    0.0062471656
    0.318907619
    0.8254657

    Empenigr
    0.1164701677
    0.137604904
    0.1953245

    Rhodtome
    0.0999089739
    0.169697909
    0.1824153

    Vaccmyrt
    0.2361482843
    0.240516323
    0.2406730

    Vaccviti
    0.1523704591
    0.156502301
    0.2110550

    Pinusylv
    0.0009244423
    0.004802076
    0.0060096
```

```
head(goodness(mods, summarize = TRUE))
```

```
Callvulg Empenigr Rhodtome Vaccmyrt Vaccviti Pinusylv 0.8254657 0.1953245 0.1824153 0.2406730 0.2110550 0.0060096
```

DIAGNOSTICS FOR CONSTRAINED ORDINATIONS | INERTIA DECOMPOSITION

inertcomp() decomposes the variance in samples or species in partial, constrained, and unconstrained components

Same two types of measure available by argument statistic

- statistic = "explained (default) gives the decomposition in terms of variance
- statistic = "distance" gives decomposition in terms of the the residual distance
- Can output as proportions of total via proportion = TRUE

```
head(inertcomp(mods, proportional = TRUE))
```

```
CCA CA
Callvulg 0.8254657 0.1745343
Empenigr 0.1953245 0.8046755
Rhodtome 0.1824153 0.8175847
Vaccmyrt 0.2406730 0.7593270
Vaccviti 0.2110550 0.7889450
Pinusylv 0.0060096 0.9939904
```

DIAGNOSTICS FOR CONSTRAINED ORDINATIONS | SPECIES-ENVIRONMENT CORRELATIONS

spenvcor() returns the species-environment correlation

The (weighted) correlation between the weighted average-based and the linear combination-based sets of site scores

A poor measure of goodness of fit. Sensitive to

- outliers (like all correlations)
- overfitting (using too many constraints)

Better models can have poorer species-environment correlations

```
spenvcor(mods)
```

```
CCA1 CCA2 CCA3 0.8554808 0.8132873 0.8792797
```

DIAGNOSTICS FOR CONSTRAINED ORDINATIONS | INTERSET CORRELATIONS

intersetcor() returns the interset correlations; the (weighted) correlation between
the weighted average-based site scores and each constraint variable

Another poor diagnostic

- correlation based
- focuses on a single constraint–axis combination at a time

Vector fitting (envfit()) or biplot scores (scores(model, display = "bp")) are
better alternatives

```
intersetcor(mods)
```

```
CCA1 CCA2 CCA3
Al 0.7356151 -0.1302334 0.4259623
P -0.3589749 -0.6110772 0.4478933
K -0.3768462 -0.1339192 0.7760441
```



RESTRICTED PERMUTATION TESTS

What is shuffled and how is of paramount importance for the test to be valid

Complete randomisation (default in **vegan**) assumes a null hypothesis where all observations are *independent*

Ecological / environmental data often aren't independent

- Temporal or spatial correlation
- · Clustering, repeated measures
- Nested sampling designs (Split-plots designs)
- Blocks
- ...

Permutation *must* give null distribution of the test statistic whilst preserving the *dependence* between observations

Trick is to shuffle the data whilst preserving that dependence

RESTRICTED PERMUTATIONS

Canoco has had restricted permutations for a *long* time. **vegan** has only recently caught up & we're not (quite) there yet

vegan still only knows how to completely randomise data or completely randomise within blocks (called **strata** in **vegan**)

The newish package **permute** grew out of initial code in the **vegan** repository to generate the sorts of restricted permutations available in Canoco

We have fully integrated **permute** into **vegan**...

vegan depends on **permute** so you should have it already installed & loaded when using **vegan**

RESTRICTED PERMUTATIONS WITH PERMUTE

permute follows Canoco closely — at the chiding of Cajo ter Braak when it didn't do what he wanted!

Samples can be thought of as belonging to three levels of a hierarchy

- the sample level; how are individual samples permuted
- the *plot* level; how are samples grouped at an intermediate level
- the block level; how are samples grouped at the outermost level

Blocks define groups of plots, each of which can contain groups of samples

RESTRICTED PERMUTATIONS WITH PERMUTE

Blocks are *never* permuted; if defined, only plots or samples *within* the blocks get shuffled & samples never swapped between blocks

Plots or samples within plots, or both can be permuted following one of four simple permutation types

- 1. Free permutation (randomisation)
- 2. Time series or linear transect, equal spacing
- 3. Spatial grid designs, equal regular spacing
- 4. Permutation of plots (groups of samples)
- 5. Fixed (no permutation)

Multiple plots per block, multiple samples per plot; plots could be arranged in a spatial grid and samples within each of the plots form a time series

RESTRICTED PERMUTATIONS WITH PERMUTE | BLOCKS

Blocks are a random factor that does not interact with factors that vary within blocks

Blocks form groups of samples that are never permuted between blocks, only within blocks

Using blocks you can achieve what the **strata** argument does now in **vegan**; needs to be a factor variable

The variation *between* blocks should be excluded from the test; **permute** doesn't do this for you!

Use + Condition(blocks) in the model formula where blocks is a factor containing the block membership for each observation

RESTRICTED PERMUTATIONS WITH PERMUTE | TIME SERIES & LINEAR TRANSECTS

Can link *randomly* starting point of one series to any time point of another series if series are stationary under null hypothesis that the series are unrelated

Achieve this via cyclic shift permutations — wrap series into a circle by joining start and end points

Works OK if there are no trends or cyclic pattern — autocorrelation structure only broken at the end points if series are stationary

Can detrend to make series stationary but not if you want to test significance of a trend

```
shuffle(10, control = how(within = Within(type = "series")))
```

```
[1] 9 10 1 2 3 4 5 6 7 8
```

RESTRICTED PERMUTATIONS WITH PERMUTE | SPATIAL GRIDS

- The trick of cyclic shifts can be extended to two dimensions for a regular spatial grid arrangement of points
 - Now shifts are toroidal as we join the end point in the x direction together and in the y direction together

```
matrix(perm, ncol = 3)
```

```
[1,1] [,2] [,3]
[1,] 6 9 3
[2,] 4 7 1
[3,] 5 8 2
```

RESTRICTED PERMUTATIONS WITH PERMUTE | WHOLE-PLOTS & SPLIT-PLOTS I

Split-plot designs are hierarchical with two levels of units

- 1. whole-plots, which contain
- 2. split-plots (the samples)

Can permute one or both of these but whole-plots must be of equal size

Essentially allows more than one error stratum to be anlyzed

Test effect of constraints that vary *between* whole plots by permuting the whole-plots whilst retaining order of split-splots (samples) within the whole-plots

Test effect of constraints that vary within whole-plots by permuting the split-plots within whole-plots without permuting the whole-plots

RESTRICTED PERMUTATIONS WITH PERMUTE | WHOLE-PLOTS & SPLIT-PLOTS II

Whole-plots or split-plots can be time series, linear transects or rectangular grids in which case the appropriate restricted permutation is used

If the split-plots are parallel time series & time is an autocorrelated error component affecting all series then the same cyclic shift can be applied to each time series (within each whole-plot) (constant = TRUE)

RESTRICTED PERMUTATIONS WITH PERMUTE | MIRRORING

(Without wanting to get *too* technical) Mirroring in restricted permutations allows for isotropy in dependences by reflecting the ordering of samples in time or spatial dimensions

For a linear transect, technically the autocorrelation at lag h is equal to that at lag -h (also in a trend-free time series)

Hence the series (1, 2, 3, 4) and (4, 3, 2, 1) are equivalent fom this point of view & we can draw permutations from either version

Similar argument can be made for spatial grids

Using mirror = TRUE then can double (time series, linear transects) or quadruple (spatial grids) the size of the set of permutations

RESTRICTED PERMUTATIONS WITH PERMUTE | THE SET OF PERMUTATIONS

Using restricted permutations can severely reduce the size of the set of allowed permutations

As the minimum p value obtaininable is 1/np where np is number of allowed permutations (including the observed) this can impact the ability to detect signal/pattern If we don't want mirroring

- in a time series of 20 samples the minimum p is 1/20 (0.05)
- in a time series of 100 samples the minimum p is 1/100 (0.01)
- in a data set with 10 time series each of 20 observations (200 total), if we assume an autocorrelated error component over all series (constant = TRUE) then there are only 20 permutations of the data and minimum p is 0.05

When the set of permutations is small it is better to switch to an exact test & evaluate all permutations in the set rather than randomly sample from the set

In permute, we set up a permutation scheme with how()

We sample from the permutation scheme with

- shuffle(), which gives a single draw from scheme, or
- \cdot shuffleSet(), which returns a set of n draws from the scheme

allPerms() can generated the entire set of permutations — note this was designed for small sets of permutations & is slow if you request it for a scheme with many thousands of permutations!

how() has three main arguments

- 1. within takes input from helper Within()
- plots takes input from helper Plots()
- 3. **blocks** takes a factor variable as input

```
plt <- gl(3, 10)
h <- how(within = Within(type = "series"), plots = Plots(strata = plt))</pre>
```

Helper functions make it easy to change one or a few aspects of permutation scheme, rest left at defaults

```
args(Within)

function (type = c("free", "series", "grid", "none"), constant = FALSE,
    mirror = FALSE, ncol = NULL, nrow = NULL)

NULL

args(Plots)

function (strata = NULL, type = c("none", "free", "series", "grid"),
    mirror = FALSE, ncol = NULL, nrow = NULL)

NULL
```

how() has additional arguments, many of which control the heuristics that kick in to stop you shooting yourself in the foot and demanding 9999 permutations when there are only 10

- complete should we enumerate the entire set of permutations?
- minperm lower bound on the size of the set of permutations at & below which we turn on complete enumeration

```
args(how)
```

RESTRICTED PERMUTATIONS WITH PERMUTE | TIME SERIES EXAMPLE I

Time series within 3 plots, 10 observation each

```
plt <- gl(3, 10)
h <- how(within = Within(type = "series"), plots = Plots(strata = plt))
set.seed(4)
p <- shuffle(30, control = h)
do.call("rbind", split(p, plt)) ## look at perms in context</pre>
[.1] [.2] [.3] [.4] [.5] [.6] [.7] [.8] [.9] [.10]
```

```
[,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10]
1 7 8 9 10 1 2 3 4 5 6
2 12 13 14 15 16 17 18 19 20 11
3 24 25 26 27 28 29 30 21 22 23
```

RESTRICTED PERMUTATIONS WITH PERMUTE | TIME SERIES EXAMPLE II

27

28

29 30 21

Time series within 3 plots, 10 observation each, same permutation within each

22 23 24 25 26

RESTRICTED PERMUTATIONS WITH PERMUTE | WORKED EXAMPLE WITH VEGAN

Here by dragons

```
## Analyse the Ohraz data Case study 5 of Leps & Smilauer

## load vegan
library("vegan")

## load the data
spp <- read.csv("data/ohraz-spp.csv", header = TRUE, row.names = 1)
env <- read.csv("data/ohraz-env.csv", header = TRUE, row.names = 1)
molinia <- spp[, 1]
spp <- spp[, -1]

## Year as numeric
env <- transform(env, year = as.numeric(as.character(year)))</pre>
```

RESTRICTED PERMUTATIONS WITH PERMUTE | WORKED EXAMPLE WITH VEGAN

```
## hypothesis 1
c1 <- rda(spp ~ year + year:mowing + year:fertilizer +
          vear:removal + Condition(plotid). data = env)
h <- how(within = Within(type = "none").</pre>
         plots = Plots(strata = env$plotid, type = "free"))
set.seed(42)
anova(c1. permutations = h. model = "reduced")
Permutation test for rda under reduced model
Plots: env$plotid, plot permutation: free
Permutation: none
Number of permutations: 199
Model: rda(formula = spp ~ year + year:mowing + year:fertilizer + year:removal + Condition(plotid), data = env)
         Df Variance
                          F Pr(>F)
         4 158 85 6 4247 0 005 **
Model
Residual 98 556 38
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
anova(c1. permutations = h. model = "reduced". bv = "axis")
```

Permutation test for rda under reduced model Marginal tests for axes Plots: env\$plotid, plot permutation: free Permutation: none Number of permutations: 199

RESTRICTED PERMUTATIONS WITH PERMUTE | WORKED EXAMPLE WITH VEGAN

```
## hypothesis 2
c2 <- rda(spp ~ year:mowing + year:fertilizer + year:removal +
          Condition(vear + plotid). data = env)
anova(c2. permutations = h. model = "reduced")
Permutation test for rda under reduced model
Plots: env$plotid, plot permutation: free
Permutation: none
Number of permutations: 199
Model: rda(formula = spp ~ year:mowing + year:fertilizer + year:removal + Condition(year + plotid), data = env)
        Df Variance
                         F Pr(>F)
              99 24 5 3517 0 005 **
Model
Residual 90 556.30
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
anova(c2. permutations = h. model = "reduced". by = "axis")
Permutation test for rda under reduced model
Marginal tests for axes
Plots: env$plotid, plot permutation: free
Permutation: none
Number of permutations: 199
Model: rda(formula = spp ~ year:mowing + year:fertilizer + year:removal + Condition(year + plotid), data = env)
        Df Variance
                          F Pr(>F)
              54 14 8 7595 A AA5 ++
RDA1
RDA2
              34.28 5.5458 0.005 **
RDA3
              10.82 1.7499 0.440
```

REFERENCES

Blanchet, F Guillaume, Pierre Legendre, and Daniel Borcard. 2008. "Forward Selection of Explanatory Variables." *Ecology* 89 (9). Eco Soc America: 2623–32.